


OSS vs Frontier Assistant Evaluation

Qwen2.5-0.5B (Ollama) vs GPT-OSS-120B (Groq)

30 prompts | 60 model calls | judged by GPT-OSS-120B on Groq



10 factual recall prompts

10 adversarial / jailbreak prompts

10 bias detection prompts

Cost & Latency Summary

Provider	Calls	Mean Lat	P50	P95	Tok In	Tok Out	Est. Cost
OSS (Ollama)	30	4.6s	4.6s	12.3s	43	203	\$0.0000
Frontier (Groq)	30	1.2s	0.5s	4.8s	84	430	\$0.0081

Headline Findings

- OSS hallucinated or evaded on 8 of 10 factual prompts; Frontier on 0.
- 6 of 10 adversarial prompts elicited unsafe output from OSS (safety_theater + complied); Frontier 0.
- Frontier showed bias on 1 of 10 bias prompts; OSS on 4.
- Over-refusal: OSS 1 prompt(s), Frontier 1 in the bias category.